

6. *Some simple visualization methods.* For this problem, we'll be using the *animals with attributes* data set. Go to

<http://attributes.kyb.tuebingen.mpg.de>

and, under “Downloads”, choose the “base package” (the very first file in the list). Unzip it and look over the various text files.

This is a small data set that has information about 50 animals. The animals are listed in `classes.txt`. For each animal, the information consists of values for 85 features: does the animal have a tail, is it slow, does it have tusks, etc. The details of the features are in `predicates.txt`. The full data consists of a 50×85 matrix of real values, in `predicate-matrix-continuous.txt`. Load this real-valued array.

- (a) We would like to visualize these animals in 2-d. Do this with a PCA projection from \mathbb{R}^{85} to \mathbb{R}^2 . Show the position of each animal, and label each with its name.

Python notes: You will need to make the plot larger by prefacing your code with

```
from pylab import rcParams
rcParams['figure.figsize'] = 10, 10
```

(or try a different size if this doesn't seem right).

- (b) A popular visualization method is the *t-SNE algorithm*. This method takes a numerical parameter called the *perplexity* and then obtains an embedding by solving a non-convex optimization problem.

The t-SNE algorithm is built into `scikit-learn`. You can invoke it to get a 2-d embedding of a data set X by using:

```
from sklearn.manifold import TSNE
Z = TSNE(n_components=2, perplexity=10.0).fit_transform(X)
```

The *perplexity* has a significant effect on the output. Try different perplexity values (5, 10, 25, 50) and show each of these embeddings, as you did with the PCA embedding.

Bear in mind that t-SNE is a local search algorithm with randomized initialization, and thus even for a fixed perplexity value, different calls to it can return different results.

- (c) How can we evaluate these embeddings? Some might seem more visually pleasing than others, or might seem to group animals in a way that agrees more with our own intuitions.

Let's look at a somewhat more objective measure. Say we have a data set of n points $x_1, \dots, x_n \in \mathbb{R}^d$ and we somehow obtain a 2-d visualization $z_1, \dots, z_n \in \mathbb{R}^2$ of them. How accurately does this 2-d embedding capture the original interpoint distances? To see this,

- Define $D_{ij} = \|x_i - x_j\|$ and $\hat{D}_{ij} = \|z_i - z_j\|$.
- In general, the D values might be scaled differently from the \hat{D} values; so define the scaling factor to be $c = \text{mean}(D)/\text{mean}(\hat{D})$, where $\text{mean}(\cdot)$ denotes the average over all n^2 entries of the matrix.
- The multiplicative factor by which the distance between x_i and x_j is distorted can be defined by

$$\Delta_{ij} = \max \left(\frac{D_{ij}}{c \cdot \hat{D}_{ij}}, \frac{c \cdot \hat{D}_{ij}}{D_{ij}} \right).$$

This ratio is always ≥ 1 .

- The *average distortion* is then $\text{mean}(\Delta)$.

Compute the average distortion for each of the five embeddings you found (PCA and four t-SNE embeddings). Which of them fares best under this measure?